

Master in Artificial Intelligence  
Supervised and Experiential Learning

---

Rule Induction from a Set of Exemplars  
(RISE)

---

Romero Antolin, Axel

December 30, 2023

# Contents

<b>1</b>	<b>Introduction</b>	<b>1</b>
<b>2</b>	<b>RISE</b>	<b>1</b>
<b>3</b>	<b>Methodology</b>	<b>2</b>
3.1	Implementation . . . . .	2
3.2	Evaluation . . . . .	5
<b>4</b>	<b>Results</b>	<b>6</b>
4.1	Toy data set . . . . .	6
4.2	Heart Disease . . . . .	7
4.3	Breast Cancer . . . . .	8
<b>5</b>	<b>Conclusions</b>	<b>9</b>

# 1 Introduction

In this first practical work of the Supervised and Experiential Learning course, we implemented and validated the Rule Induction from a Set of Exemplars (RISE) rule-based classifier. A rule-based classifier is a type of classifier that makes predictions based on a set of rules. These rules can be simple if-then statements or more complex logical expressions. The classifier evaluates each rule and applies the corresponding prediction when the rule conditions are met. Section 2 provides a detailed explanation of the RISE algorithm, which is a method for automatic learning of the rules given a data set.

The implementation was done in python and the code is available in the `source` folder. Further details about the organization of the code and how to execute it are given in Section 3. For the validation of the implementation of the RISE model, we used three data sets of different sizes to measure the performance and ability to classify new instances. The results of the tests can be seen in Section 4.

## 2 RISE

The RISE algorithm is a type of rule-based classifier first implemented by Pedro Domingos [1] where the rules are learned by gradually generalizing instances until there is no improvement in accuracy. The main characteristics of this algorithm are that it can handle numerical variables, which is something not very usual in other famous algorithms like RULES, PRISM and CN2, and that it produces a high number of rules that are, in general, very specific. This happens because initially, all the instances of the training data are treated as a rule and they are generalised in order to improve the performance.

The rules in RISE are composed of a consequent part (predicted class) and an antecedent part, which is a combination of conditions. Each condition involves only one attribute whereas for categorical variables the condition is an equality test and for numeric variables, the condition is an interval closed on both sides.

The first part of the algorithm is learning the rules. First, all the instances of the data set are treated as rules and we calculate the initial accuracy of the models using leave-one-out. For each instance, we remove the instance and the related rule and use the rest of them to classify the instance and finally obtain an average performance of the set of rules.

Then, for each rule in the set of rules, we find the nearest instance to the rule not covered by the rule and being of the same class. We said that an instance is covered by a rule if all the attributes of the instance satisfy the antecedent of the rule. The distance  $\Delta(R, E)$  used in this algorithm between  $R$  and  $E$ , where  $R = (A_1, A_2, \dots, A_p, C_R)$  with class  $C_R$  and condition  $A_i$  for the  $i$ th attribute and  $E = (e_i, e_2, \dots, e_p, C_E)$  with class  $C_E$  and value  $e_i$  for the  $i$ th attribute, is defined as:

$$\Delta(R, E) = \sum_{i=1}^a \delta^s(i) \quad (1)$$

where  $s$  is a natural-valued parameter ( $s = 1, 2, 3, \dots$ ), and the component distance  $\delta(i)$  for the  $i$ th attribute is:

$$\delta(i) = \begin{cases} 0 & \text{if } A_i = \text{True} \\ SVDM(r_i, e_i) & \text{if } i \text{ is categorical and } A_i \neq \text{True} \\ \delta_{num}(i) & \text{if } i \text{ is numeric and } A_i \neq \text{True} \end{cases} \quad (2)$$

where  $SVDM(r_i, e_i)$  is the simplified value difference metric defined as:

$$SVDM(x_i, x_j) = \sum_{h=1}^c |P(C_h|x_i) - P(C_h|x_j)|^q \quad (3)$$

where  $c$  is the number of classes,  $C_h$  is the  $h$ th class, and  $q$  is a natural-valued parameter. Then, the distance metric for the numerical attributes is defined as:

$$\delta_{num}(i) = \begin{cases} 0 & \text{if } r_{i,lower} \leq e_i \leq r_{i,upper} \\ \frac{e_i - r_{i,upper}}{e_{i,max} - e_{i,min}} & \text{if } e_i > r_{i,upper} \\ \frac{r_{i,lower} - e_i}{e_{i,max} - e_{i,min}} & \text{if } e_i < r_{i,lower} \end{cases} \quad (4)$$

Once we have the nearest instance to the actual rule, we generalize the rule in order to make it cover the instance. We have the rule  $R = (A_1, A_2, \dots, A_p, C_R)$  and the instance  $E = (e_1, e_2, \dots, e_p, C_E)$  where the condition  $A_i$  is either *True*,  $e_i = r_i$  (categorical) or  $r_{i,lower} \leq e_i \leq r_{i,upper}$ . For each attribute  $i$ , if  $A_i = \text{True}$  then do nothing; if  $i$  is categorical and  $e_i \neq r_i$  then  $A_i = \text{True}$ ; if  $i$  is numerical and  $e_i > r_{i,upper}$  then  $r_{i,upper} = e_i$ ; and if  $i$  is numeric and  $e_i < r_{i,lower}$  then  $r_{i,lower} = e_i$ .

With the generalized rule, we create a new set of rules where we replace the actual rule with the generalized one. With this new set of rules, we calculate the accuracy using the same procedure of leave-one-out and, if the value is higher than before we maintain this new set of rules.

We repeat this process until there is no improvement in the accuracy of the set of rules. When the process has finished, we have the final set of rules derived from the training data that will be used for the prediction of new instances, which is the second part of the model. In this algorithm, we classify new instances by assigning them the class of the nearest rule using the specific distance metric defined before.

### 3 Methodology

In this section, we explain the implementation of the RISE algorithm in Python 3 and the experimental approach to analyze the performance of the algorithm and the ability to produce good classifications.

#### 3.1 Implementation

In the previous section, we explained the theoretical background of the RISE algorithm that we used to implement it in Python 3 following the original paper [1] with some small changes. All the code is available in the `code` folder and the instructions to execute the code will be given at the end of the section. We created a main class called `RISE` with the following parameters:

- **dataset\_name**: name of the data set to use to train and evaluate the RISE algorithm. This value has to be one valid name of a csv file in the `data` folder.
- **target\_name**: name of the class variable.
- **cat\_variables**: list with the name of the categorical variables of the data set.
- **test\_split**: the value of the percentage of instances of the data set separated for the evaluation phase, the test data. The rest of the instances are used for training.
- **max\_iter**: maximum number of iterations in the training process.

There are different important parts in the class of the algorithm that we are going to explain in detail in the following paragraphs. First, we integrated into the class the import and preprocessing of the data. We import the data set using the **dataset\_name** parameter where the file has to be in csv format. For the preprocessing, first, we check if there are missing data and, in the case of missing values, we imputed them with mean or mode. Then, for the categorical variables of the data set, we encoded the values in order to avoid characters and have integers that make it easier to calculate distances and manage the rules. Finally, we performed the train and test split of the data for the training and evaluation using the test size of the parameter **test\_split**.

Then, to represent the rules and be able to calculate distances between them and instances, we

use arrays where we store the values of the attributes to form the antecedent part. For the categorical variables, we have the value and for the numerical variables we have an interval represented as “value-value”. The last value of the array is the consequent part, the predicted class.

The next step is the implementation of the training phase with the `fit()` method. In Algorithm 1 we can see the implemented algorithm of the main procedure of the RISE that we explained in Section 2. The initial part is the initialization of the set of rules where each instance is transformed into a rule with the `initialize_rules()` method. As a result, we obtain an array with all the rules where the size is equal to the size of the training data. Then, we calculate the initial accuracy of the set of rules with `estimate_accuracy()` and start the while loop until there is no improvement in the accuracy. For the estimation of the accuracy, we used the leave-one-out method where, when we predict the class of an instance, we remove the rule related to that instance and try to predict it with the rest of the rules. If we do not add this procedure we will obtain always a precision of 1 since for the prediction of each rule we will use the rule created from that instance. The method `estimate_accuracy()` classifies each instance with the class label of the nearest rule using the distance function explained in Section 2. The accuracy can be estimated with the actual set of rules or with a specific set that can be introduced as a parameter.

Then, for each rule, we find the nearest instance to the rule not covered and with the same class label with the `nearest_instance_to_rule()` and the process of generalization of the rule, explained in Section 2, is implemented in the `rule_generalization()`. Finally, we evaluate the new accuracy of the updated set of rules and decide whether keep the new rule or not. After some steps, we obtained the final set of rules that can be used to classify new instances. In this implementation, we add a maximum number of iterations, which is something that was not implemented in the original paper.

We also add a procedure for the elimination of repeated rules that was not explained in the original pseudo-code. In Algorithm 1 this procedure is also added. It consists of checking if each rule is also present in the previous rules and, if not, continuing with the normal procedure. After the generalization of the rule, we check again if the created rule is already present in the set of rules, if not we keep the rule.

Once we have the final set of rules derived from the training data, we implemented a method for the prediction of new instances using the trained set of rules, `predict()`. Also, to evaluate the performance of the set of rules, we implemented the `evaluate()` method which, for both train and test, computes different metrics to evaluate the model and creates a txt file with the value of the metrics. The details of the evaluation metrics will be given in the next section.

For the execution of the RISE algorithm and use it with a custom data set we have to use the RISE class and the `main.py` script in the `source` folder containing the code used in the experimental analysis. To use a custom data set we have to add it to the data folder in a csv file. To train the model we have to execute the `fit()` method that returns the final set of rules and the `evaluation()` method to test the model and obtain a txt file with the metrics.

During the implementation of the algorithm we had to face different problems, but the most important and the one that we dedicate more time was the optimization of the training. The execution time of the training is very long even with the process of optimization and this is due to the high number of times that we have to iterate over all the instances of the data set. Because of this problem, we could not manage to test the model with a data set with a high number of instances and it is the main problem of this algorithm.

---

**Algorithm 1** RISE algorithm

---

```
rule_set  $\leftarrow$  initialize_rules()  
initial_accuracy  $\leftarrow$  estimate_accuracy(Rule_set)  
acc  $\leftarrow$  Initial_accuracy  
acc_increased  $\leftarrow$  True  
c  $\leftarrow$  0  
while acc_increased and c < max_iter do  
  for rule in Rule_set do  
    new_set_rules  $\leftarrow$  copy(rule_set)  
    if rule repeated then  
      new_set_rules[i]  $\leftarrow$  None  
    else  
      nearest_instance  $\leftarrow$  nearest_instance_to_rule(rule)  
      new_rule  $\leftarrow$  rule_generalization(rule, nearest_instance)  
      new_set_rules[i]  $\leftarrow$  new_rule  
    end if  
    new_acc  $\leftarrow$  estimate_accuracy(new_set_rules)  
  end for  
  if new_acc < acc then  
    acc  $\leftarrow$  new_acc  
    if new_rule repeated then  
      new_set_rules[i]  $\leftarrow$  None  
    end if  
    set_rules  $\leftarrow$  new_set_rules  
  end if  
  acc_increased  $\leftarrow$  acc > initial_accuracy  
  initial_accuracy  $\leftarrow$  acc  
  c  $\leftarrow$  c + 1  
end while
```

---

### 3.2 Evaluation

For the evaluation of the implementation of the RISE algorithm, we tested it in three data sets from the UCI Machine Learning Repository of different sizes in order to measure the ability to predict new instances. As mentioned in the previous section, the execution time of the RISE algorithm is very long and we could not manage to test it in a large data set. We decided to, a part of the UCI data sets, include one simple and small data set from the theoretical slides used to test the implementation and it is easy to interpret the results. The data sets are:

- **Heart Disease.** Data set with 270 instances and 14 attributes where the class variable is **disease**. The set of variables contains numerical and categorical attributes and there is no missing data. The variables are:
  - age: age in years
  - sex: sex (1 = male; 0 = female)
  - chest: chest pain type where 1: typical angina; 2: atypical angina; 3: non-anginal pain; and 4: asymptomatic
  - rbp: resting blood pressure (in mm Hg on admission to the hospital)
  - cholesterol: serum cholesterol in mg/dl
  - sugar: (fasting blood sugar > 120 mg/dl) (1 = true; 0 = false)
  - ecg: resting electrocardiographic results where 0: normal; 1: having ST-T wave abnormality (T wave inversions and/or ST elevation or depression of > 0.05 mV); and 2: showing probable or definite left ventricular hypertrophy
  - max hr: maximum heart rate achieved
  - angina: exercise-induced angina (1 = yes; 0 = no)
  - oldpeak: ST depression induced by exercise relative to rest
  - slope: the slope of the peak exercise ST segment where 1: upsloping; 2: flat; and 3: downsloping
  - vessels: number of major vessels (0-3) coloured by fluoroscopy
  - thal: 3 = normal; 6 = fixed defect; 7 = reversible defect
  - disease: diagnosis of heart disease (angiographic disease status) where 0: < 50% diameter narrowing; and 1: > 50% diameter narrowing
- **Breast Cancer.** Data set with 699 instances and 10 attributes where the class variable is **Class**. In this case, the set of variables contains only numerical attributes and there are 16 instances that have missing data. The variables are:
  - Clump Thickness: Represents the thickness of the tumour in mm
  - Uniformity of Cell Size: Represents the uniformity in the size of the cells
  - Uniformity of Cell Shape: Represents the uniformity in the shape of the cells
  - Marginal Adhesion: Represents the degree of adhesion of the tumour cells to nearby cells
  - Single Epithelial Cell Size: Represents the size of the epithelial cells. It is a numerical attribute
  - Bare Nuclei: Represents the presence of a nucleus in the tumour cells
  - Bland Chromatin: Represents the degree of chromatin staining in the tumour cells
  - Normal Nucleoli: Represents the size and shape of the nucleoli in the tumour cells
  - Mitoses: Represents the number of mitoses (cell division) in the tumour cells
  - Class: Represents the diagnosis of the tumour (2 for benign, 4 for malignant)

In order to measure the performance and ability to classify new instances of the RISE model we used

train/test cross-validation where the training split was used to train the model and the test split to evaluate the performance in a new set of instances. The metrics used for the evaluation are:

- **Accuracy:** Train and test general accuracy of the model. This is the ratio between the number of instances correctly classified using the set of rules and the total number of instances.
- **Rule recall:** The ratio between the number of instances satisfying the antecedent of the rule and the consequent of the rule, and the total number of instances.
- **Rule accuracy:** The ratio between the number of instances satisfying the antecedent of the rule and the consequent of the rule, and the number of instances satisfying the antecedent of the rule.
- **Rule coverage:** The ratio between the number of instances satisfying the antecedent of the rule and the total number of instances.

## 4 Results

In this section, we will present the results obtained in the evaluation of the RISE algorithm in the three data sets and the procedure explained in the previous section. We will present the values of the metrics and the most important rules derived from the data for each one of the data sets. In the **results** folders there are the generated files with the detailed description of the results for the different tests, like all the rules generated and their accuracy, recall and coverage.

### 4.1 Toy data set

This small data set was used for testing the implementation of the RISE algorithm in order to find bugs. As it is a small data set with a small number of attributes it is easy to interpret. The general accuracy for both training and testing is 1.

The final set of rules consists of 18 rules using the training data which has 19 instances. We can see that more or less each instance is transformed into a final rule showing the high amount of rules that this algorithm generates. For example, the first rule says that a young person with astigmatism and normal tear production is classified with a hard label. The recall of the rules is 0.67 which indicates that for the cases with the label `RecomendedLense = hard`, 66.7% are covered with the rule. There are also rules for the other classes of the predicted attribute.



Rule	Coverage	Accuracy	Recall
(Age = young) AND (Astigmatism = yes) AND (TearProduction = normal) THEN Recomendense = hard	0.105	1	0.667
(VisualDeficiency = myopia) AND (Astigmatism = yes) AND (TearProduction = normal) THEN Recomendense = hard	0.105	1	0.667
(VisualDeficiency = hypermetropia) AND (Astigmatism = no) AND (TearProduction = normal) THEN Recomendense = soft	0.157	1	0.6
(VisualDeficiency = hypermetropia) AND (Astigmatism = no) THEN Recomendense = none	0.263	1	0.454
⋮	⋮	⋮	⋮

Table 1: Most important rules for the toy data set.

## 4.2 Heart Disease

The training of the rules for the Heart Disease data set was done with 80% of the data, which is 216 instances, and the rest, 54 instances, were used to evaluate the model. The training took approximately 2h and a total of 8 iterations (the maximum number of iterations was 10).

After the training, we obtain a total accuracy in the training set of 0.990 and a value of 0.769 in the test set. This shows that the RISE algorithm, for this data set, is capable of classifying the instances between the different classes with high precision. The final set of rules consists of 122 different rules where we can see some examples of the most important rules obtained from the data with the RISE algorithm in Table 2 together with the accuracy, recall and coverage of each one. For all the rules we obtain a very low coverage, which indicates that the rules are specific and they only cover a few instances. The high accuracy of the rules is also explained by the low coverage which means that the small number of instances covered by the rule are of the same class as the rule. Then, if we analyze the recall of the rules we can see that, with the three first rules, we cover the 33% of the instances of the class disease = 0.

In the `results` folder there is available the evaluation of the rules in the test set, where we can see that not all of the rules cover instances in the test. Most of the rules have a 0 value for the coverage which makes also a 0 for the recall and impossible to estimate the accuracy. This means that there are several rules that are very specific to the training data and are not very useful in the inference of new data. For the rules that have coverage greater than 0, we obtain an accuracy estimation of 1.

Rule	Coverage	Accuracy	Recall
( $52 \leq \text{age} \leq 67$ ) AND ( $115 \leq \text{rbp} \leq 152$ ) AND ( $232 \leq \text{cholesterol} \leq 564$ ) AND ( $147 \leq \text{max hr} \leq 178$ ) AND ( $\text{angina} = 0$ ) AND ( $0.4 \leq \text{oldpeak} \leq 2.3$ ) AND ( $0 \leq \text{vessels} \leq 0$ ) $\implies \text{disease} = 0$	0.060	1	0.111
( $41 \leq \text{age} \leq 59$ ) AND ( $\text{sex} = 1$ ) AND ( $118 \leq \text{rbp} \leq 135$ ) AND ( $186 \leq \text{cholesterol} \leq 258$ ) AND ( $132 \leq \text{max hr} \leq 190$ ) AND ( $\text{angina} = 0$ ) AND ( $0 \leq \text{oldpeak} \leq 1.2$ ) AND ( $0 \leq \text{vessels} \leq 0$ ) $\implies \text{disease} = 0$	0.055	1	0.102
( $29 \leq \text{age} \leq 48$ ) AND ( $\text{sex} = 1$ ) AND ( $118 \leq \text{rbp} \leq 140$ ) AND ( $182 \leq \text{cholesterol} \leq 321$ ) AND ( $\text{sugar} = 0$ ) AND ( $156 \leq \text{max hr} \leq 202$ ) AND ( $\text{angina} = 0$ ) AND ( $0 \leq \text{oldpeak} \leq 0$ ) AND ( $\text{slope} = 0$ ) AND ( $0 \leq \text{vessels} \leq 0$ ) AND ( $\text{thal} = 0$ ) $\implies \text{disease} = 0$	0.055	1	0.102
( $54 \leq \text{age} \leq 61$ ) AND ( $\text{sex} = 1$ ) AND ( $\text{chest} = 3$ ) AND ( $110 \leq \text{rbp} \leq 160$ ) AND ( $239 \leq \text{cholesterol} \leq 353$ ) AND ( $\text{sugar} = 0$ ) AND ( $88 \leq \text{max hr} \leq 145$ ) AND ( $\text{angina} = 1$ ) AND ( $0.8 \leq \text{oldpeak} \leq 3.6$ ) AND ( $\text{slope} = 1$ ) AND ( $1 \leq \text{vessels} \leq 2$ ) AND ( $\text{thal} = 2$ ) $\implies \text{disease} = 1$	0.046	1	0.101
$\vdots$	$\vdots$	$\vdots$	$\vdots$

Table 2: Most important rules for the Heart Disease data set.

### 4.3 Breast Cancer

The training of the rules for the Breast Cancer data set was done with 80% of the data, which is 559 instances, and the rest, 140 instances, were used to evaluate the model. The training took approximately 8h and a total of 4 iterations (the maximum number of iterations was 5).

The general accuracy of the final model for the training set is 0.998 and for the test set 0.964. The performance of the RISE algorithm for both training and testing is very good and we can see that the rules learned by the algorithm generalize well and are capable to make good predictions in new external data. The final set of rules consists of 273 different rules and in Table 3 we can see the top 4 most important rules based on the recall and coverage. In the **results** folder there is a csv file with all the rules ordered by higher recall and coverage and we can see that all the initial rules (the most important ones) are for the class  $\text{Class} = 0$ . Then, for all the rules in the table, we obtain a similar coverage of around 0.15 which means that each rule covers approximately 15% of the instances of the training set. For the recall, we obtain also similar values between 0.21 and 0.27 where, for example, in the first rule, the value of 0.27 means that 27% of the instances of  $\text{Class} = 0$  are covered by the rule. Then, the accuracy of the rules is one as in the case of the previous data set.

In the **results** folder it is available the full list of rules and the metrics in the training and testing data where the coverage of the rules is lower compared to the values in the train set. Also, there are two files that describe the training process with the steps and accuracies for each one and the evaluation file with the metrics and rules generated by the algorithm.

Rule	Coverage	Accuracy	Recall
$(2 \leq \text{Clump Thickness} \leq 6) \text{ AND } (1 \leq \text{Uniformity of Cell Size} \leq 2) \text{ AND } (1 \leq \text{Uniformity of Cell Shape} \leq 2) \text{ AND } (1 \leq \text{Marginal Adhesion} \leq 1) \text{ AND } (1 \leq \text{Single Epithelial Cell Size} \leq 2) \text{ AND } (1 \leq \text{Bare Nuclei} \leq 2) \text{ AND } (2 \leq \text{Bland Chromatin} \leq 7) \text{ AND } (1 \leq \text{Normal Nucleoli} \leq 1) \text{ AND } (1 \leq \text{Mitoses} \leq 1) \implies \text{Class} = 0$	0.175	1	0.270
$(1 \leq \text{Clump Thickness} \leq 2) \text{ AND } (1 \leq \text{Uniformity of Cell Size} \leq 1) \text{ AND } (1 \leq \text{Uniformity of Cell Shape} \leq 1) \text{ AND } (1 \leq \text{Marginal Adhesion} \leq 3) \text{ AND } (2 \leq \text{Single Epithelial Cell Size} \leq 2) \text{ AND } (1 \leq \text{Bare Nuclei} \leq 1) \text{ AND } (1 \leq \text{Bland Chromatin} \leq 3) \text{ AND } (1 \leq \text{Normal Nucleoli} \leq 1) \text{ AND } (1 \leq \text{Mitoses} \leq 1) \implies \text{Class} = 0$	0.152	1	0.234
$(1 \leq \text{Clump Thickness} \leq 2) \text{ AND } (1 \leq \text{Uniformity of Cell Size} \leq 1) \text{ AND } (1 \leq \text{Uniformity of Cell Shape} \leq 1) \text{ AND } (1 \leq \text{Marginal Adhesion} \leq 2) \text{ AND } (2 \leq \text{Single Epithelial Cell Size} \leq 2) \text{ AND } (1 \leq \text{Bare Nuclei} \leq 1) \text{ AND } (1 \leq \text{Bland Chromatin} \leq 3) \text{ AND } (1 \leq \text{Normal Nucleoli} \leq 1) \text{ AND } (1 \leq \text{Mitoses} \leq 1) \implies \text{Class} = 0$	0.150	1	0.231
$(1 \leq \text{Clump Thickness} \leq 3) \text{ AND } (1 \leq \text{Uniformity of Cell Size} \leq 1) \text{ AND } (1 \leq \text{Uniformity of Cell Shape} \leq 1) \text{ AND } (1 \leq \text{Marginal Adhesion} \leq 1) \text{ AND } (2 \leq \text{Single Epithelial Cell Size} \leq 2) \text{ AND } (1 \leq \text{Bare Nuclei} \leq 1) \text{ AND } (2 \leq \text{Bland Chromatin} \leq 3) \text{ AND } (1 \leq \text{Normal Nucleoli} \leq 2) \text{ AND } (1 \leq \text{Mitoses} \leq 1) \implies \text{Class} = 0$	0.136	1	0.210
⋮	⋮	⋮	⋮

Table 3: Most important rules for the Breast Cancer data set.

## 5 Conclusions

In this practical work, we implemented from scratch a rule-based classifier and evaluate its performance using different data sets. Based on the results obtained with the different data sets we can see that the classification capacity of the RISE algorithm is very high since the train and test accuracy is in general high. With this algorithm, we can create a set of rules that can be used to predict new instances with high precision. On the other hand, this algorithm presents different drawbacks that difficult its use. First, the number of generated rules is very high since at the beginning, the algorithm creates one rule for each instance and then generalizes them. In our experiments we did not use any data set that was

very big and, even that, the number of rules was very high creating more than one hundred for each data set. This makes it harder to interpret the model and the set of rules and most of the rules are very specific, which means that they are derived from the training data but are not used to predict new instances.

Another important problem of this algorithm is the execution time of the training phase. The process of checking all the rules and generalising them is very computationally expensive which makes the training time very long. For the experiments of this practical work, we have to train our model for more than 6h in order to obtain a final set of rules. If we want to use a big data set (more than 2000 instances) this algorithm is not feasible because the execution time will be very long, it will take days. Most of the time of the implementation was to try and test methods to improve the efficiency of the code and reduce the execution time so, as future work, it would be good to try to optimize more the code and have a final version of the algorithm that could handle large data sets.

## References

- [1] Pedro Domingos. “Unifying Instance-Based and Rule-Based Induction”. In: *Mach. Learn.* 24.2 (Aug. 1996), pp. 141–168. ISSN: 0885-6125. DOI: [10.1023/A:1018006431188](https://doi.org/10.1023/A:1018006431188). URL: <https://doi.org/10.1023/A:1018006431188>.